

# Lecture 2: Safety & Proof

# Policies everywhere

Carnegie Mellon University

S3 Admin Console


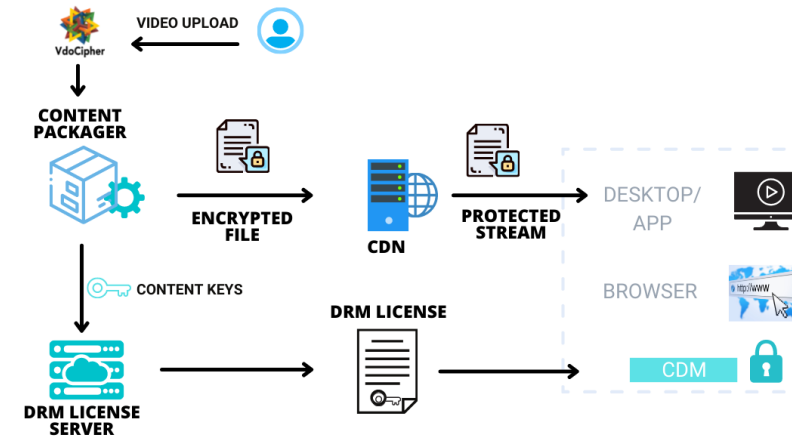
Student Course

Matthew Fredrikson

PREFERRED NAME	USER ID	PRONOUNS	PRONUNCIATION	UNIVERSITY HOLD	DIRECTORY RELEASE
mfredrik				No	Yes

Summary Academic Records Memos Documents

**Action Prevented**  
Not Authorized for this Student Summary activity.  
Resource Description: View student summary  
Access restricted because the following qualifiers failed:  
1. Evaluates whether the user is a Primary or Secondary advisor for the student. (ref: IsMyMajorAdvisee).  
2. Evaluates whether the user is an academic advisor for the student. (ref: IsMyAdvisee).  
If you feel this is in error, click [here](#) to send feedback.  
To review your Groups and Qualifiers, click 'OK' and click on 'Permissions' in the upper-right corner of S3.



# Today's goals

- Introduce safety policies
- Dive deeper into memory safety
- Understand limitations of heuristic defenses
- Formulate a good mitigation

# Safety

A class of policies that can *always* be enforced at runtime

Equivalently, policies that prohibit irremediable events

Enforcement means: possible to ensure that unsafe things **never** happen

*If the presented student ID is expired, deny entry*


*Only the student's primary advisor may view their summary*

*File's owner can read, write, execute, and others can only execute*

*Allocated memory must eventually be freed*

*Allocated memory must be freed within 500 cycles*

# Memory safety

- Goal is to prevent issues like the following:
    - Buffer overflow, over-read
    - Array index out-of-bounds
    - Uninitialized read/dereference, null dereference
    - Invalid page fault
    - Use after free, invalid/double free
- 
- Policy:** program shouldn't access illegal memory regions
- In many cases, violations result in corrupted state, unstable behavior
  - Sometimes they result in exploitable vulnerabilities
  - Why is this *safety*?

# Recap: Buffer overflow

Occurs when data is written to a location outside of the space allocated for a buffer

A buffer may be allocated:

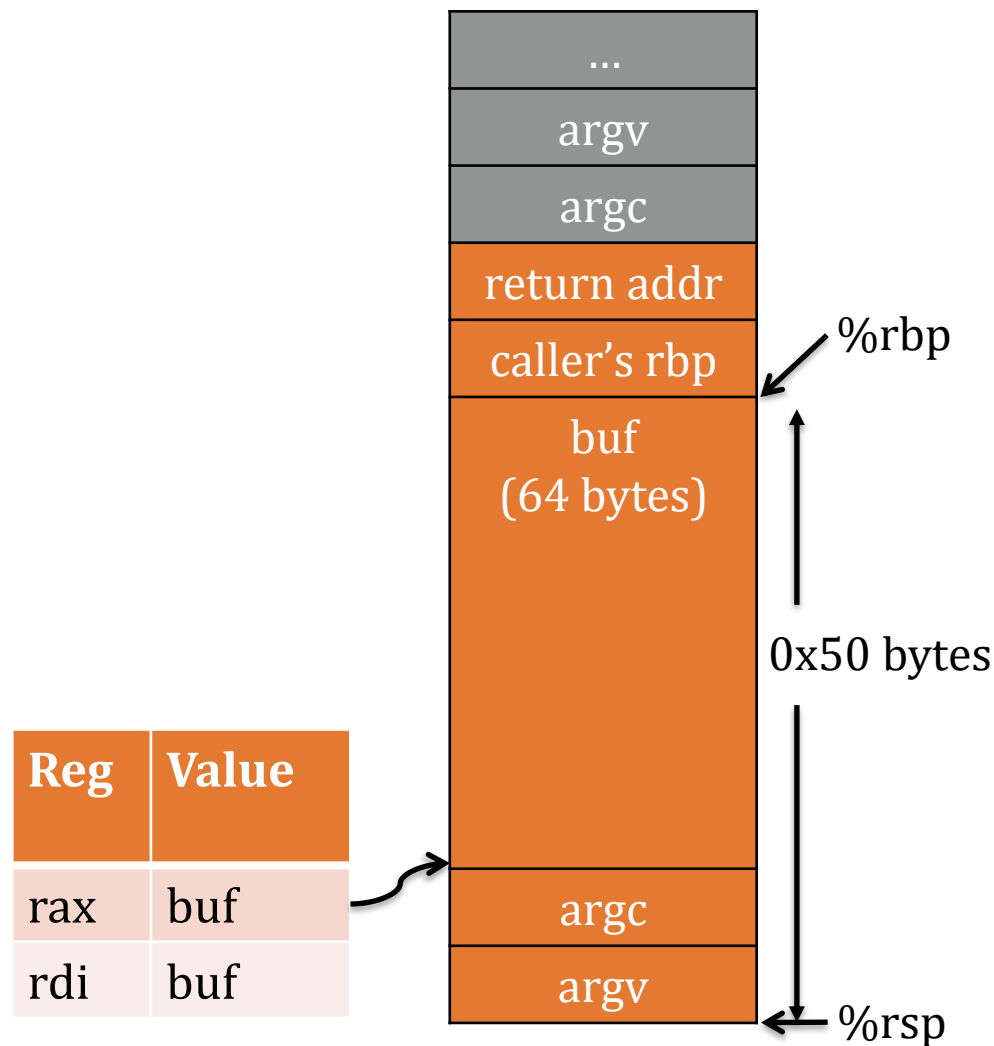
- On the stack
  - Covered today
  - Easiest case to exploit
- On the heap
  - Not covered
  - May still be exploitable, but more advanced

# Basic Example

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```




# “123456”

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```



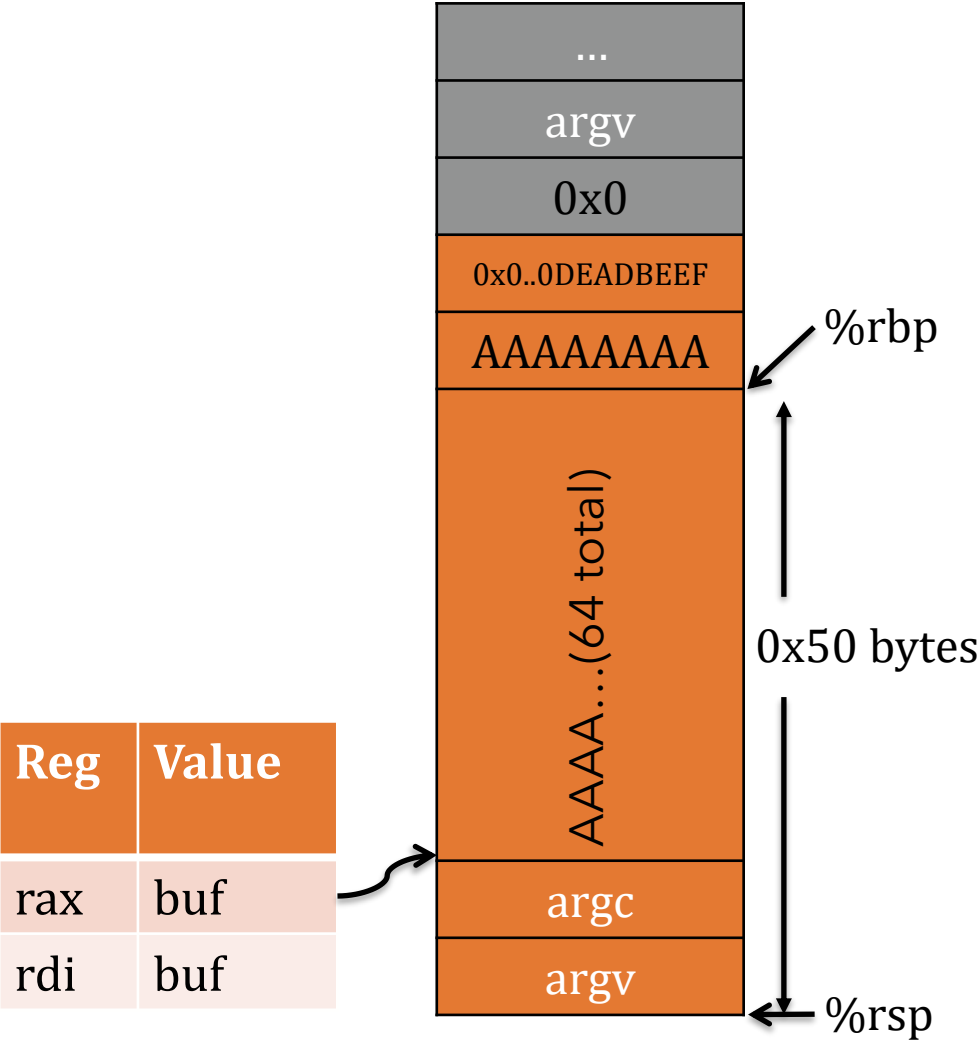


“A”x72 + “\xEF\xBE\xAD\xDE\x00\x00\x00\x00”

```
int main(int argc, char **argv) {
    char buf[64];
    gets(buf);
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp
4004fe: mov     %rsp,%rbp
400501: sub     $0x50,%rsp
400505: mov     %rdi,-0x48(%rbp)
400508: mov     %rsi,-0x50(%rbp)
40050c: lea     -0x40(%rbp),%rax
400510: mov     %rax,%rdi
400518: callq   400400 <gets@plt>
40051d: leaveq
40051e: retq
```




# Frame teardown, part 1

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```

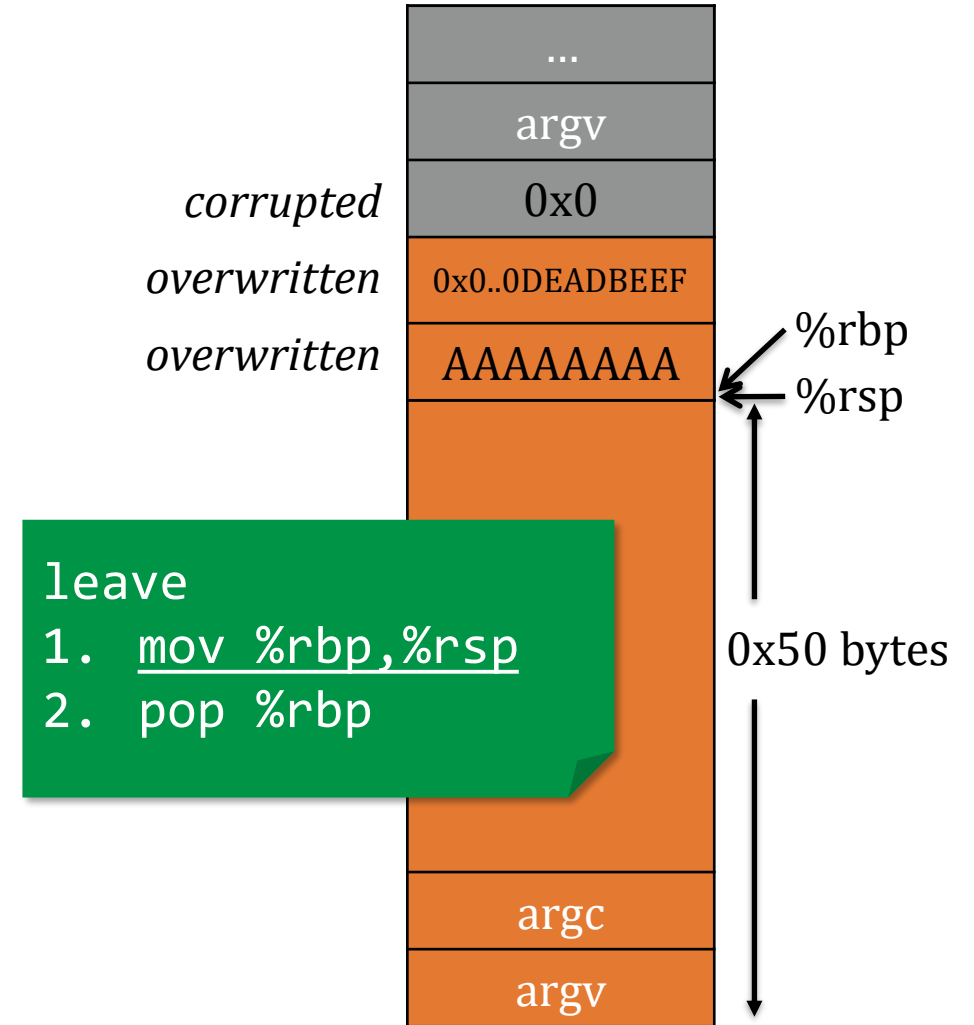


# Frame teardown, part 1

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```

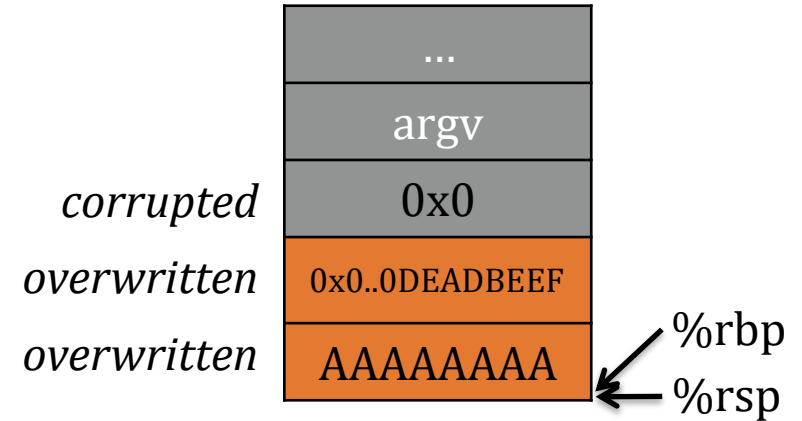


# Frame teardown, part 1

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```



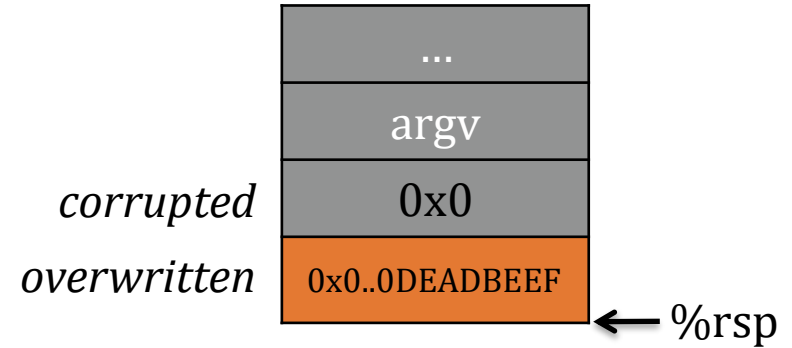
```
leave  
1. mov %rbp,%rsp  
2. pop %rbp
```

# Frame teardown, part 2

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```



%rbp = AAAAAAAAAA

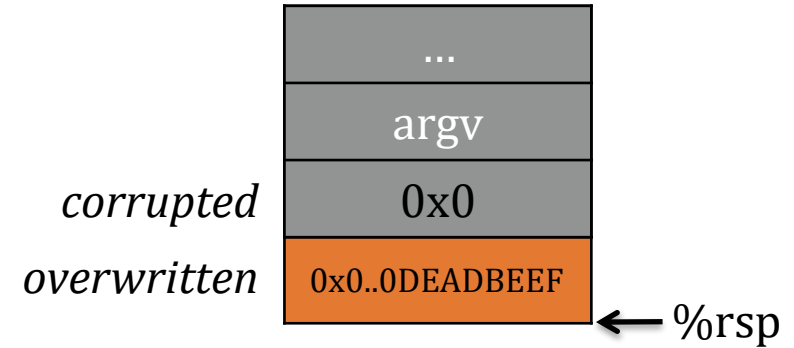
leave  
1. mov %rbp,%rsp  
2. pop %rbp

# Frame teardown, part 3

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```

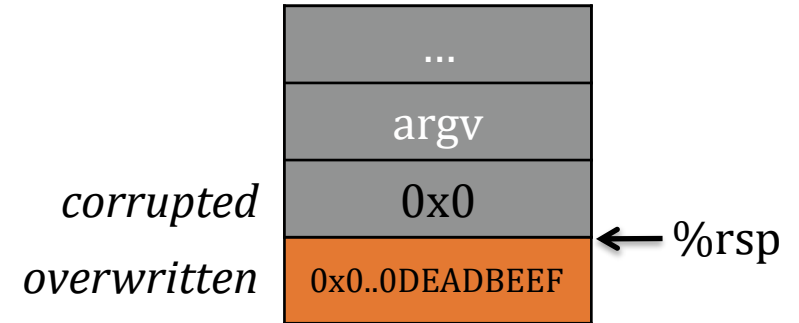


# Frame teardown, part 3

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```




# Frame teardown, part 3

```
int main(int argc, char **argv) {  
    char buf[64];  
    gets(buf);  
}
```

Dump of assembly code for function main:

```
4004fd: push    %rbp  
4004fe: mov     %rsp,%rbp  
400501: sub     $0x50,%rsp  
400505: mov     %rdi,-0x48(%rbp)  
400508: mov     %rsi,-0x50(%rbp)  
40050c: lea     -0x40(%rbp),%rax  
400510: mov     %rax,%rdi  
400518: callq   400400 <gets@plt>  
40051d: leaveq  
40051e: retq
```



**%rip = 0x00000000DEADBEEF**  
*(probably crash)*




# Shellcode

Traditionally, we inject assembly instructions for `exec("/bin/sh")` into buffer.

- see *"Smashing the stack for fun and profit"* for exact string
- or search online

```
...  
0x080483fa <+22>: call    0x8048300 <gets@plt>  
0x080483ff <+27>: leave  
0x08048400 <+28>: ret
```

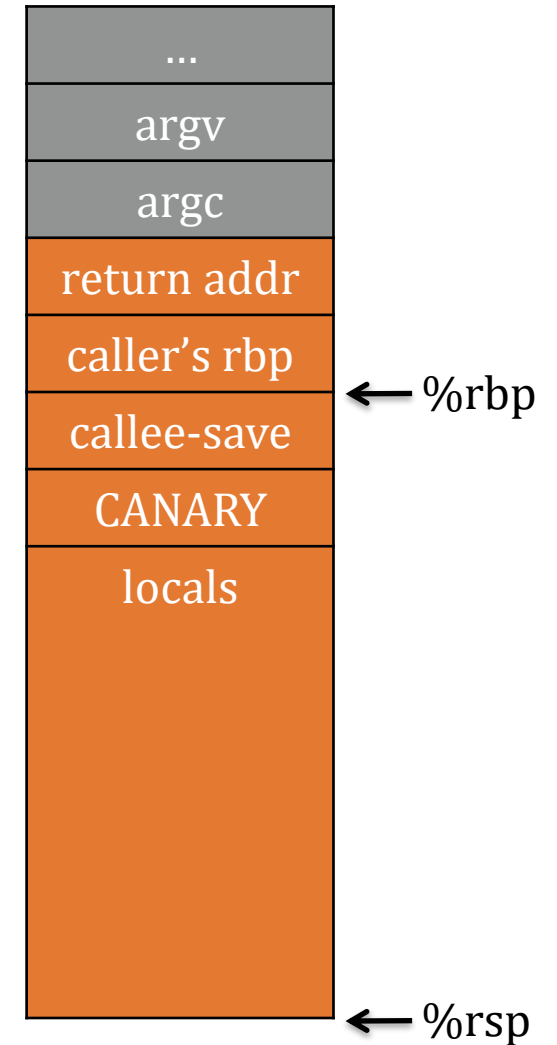


# Defenses

# StackGuard [Cowen et al. 1998]


**Key idea:** insert a unique identifier ("canary") just below the saved return address

- Function prologue inserts a random value below callee-save registers, above locals
- Epilogue checks the value before returning
- Terminate if the canary doesn't match!



# Data Pointer Subterfuge


```
int *ptr;  
char buf[64];  
memcpy(buf, user1);  
*ptr = user2;
```



# Data Pointer Subterfuge

```
int *ptr;  
char buf[64];  
memcpy(buf, user1);  
*ptr = user2;
```

First overwrite *ptr* to point  
to saved return address




# Data Pointer Subterfuge

```
int *ptr;  
char buf[64];  
memcpy(buf, user1);  
*ptr = user2;
```

First overwrite *ptr* to point  
to saved return address

Then modify return  
address via *ptr*



# Memory protection



Either (or both):

- Mark stack as non-executable
- Make sure that each page is **writable** or **executable**, exclusively

# Memory protection



CRASH

Either (or both):

- Mark stack as non-executable
- Make sure that each page is **writeable** or **executable**, exclusively




# Bypassing memory protection

Set return address to code that's already marked executable!

Fertile ground: libc

- More difficult with register-based calling conventions
- Typically done via register-loading "gadgets"

Main point: no injected code



# Layout randomization (ASLR)

**Recall:** Our exploit needs a concrete address

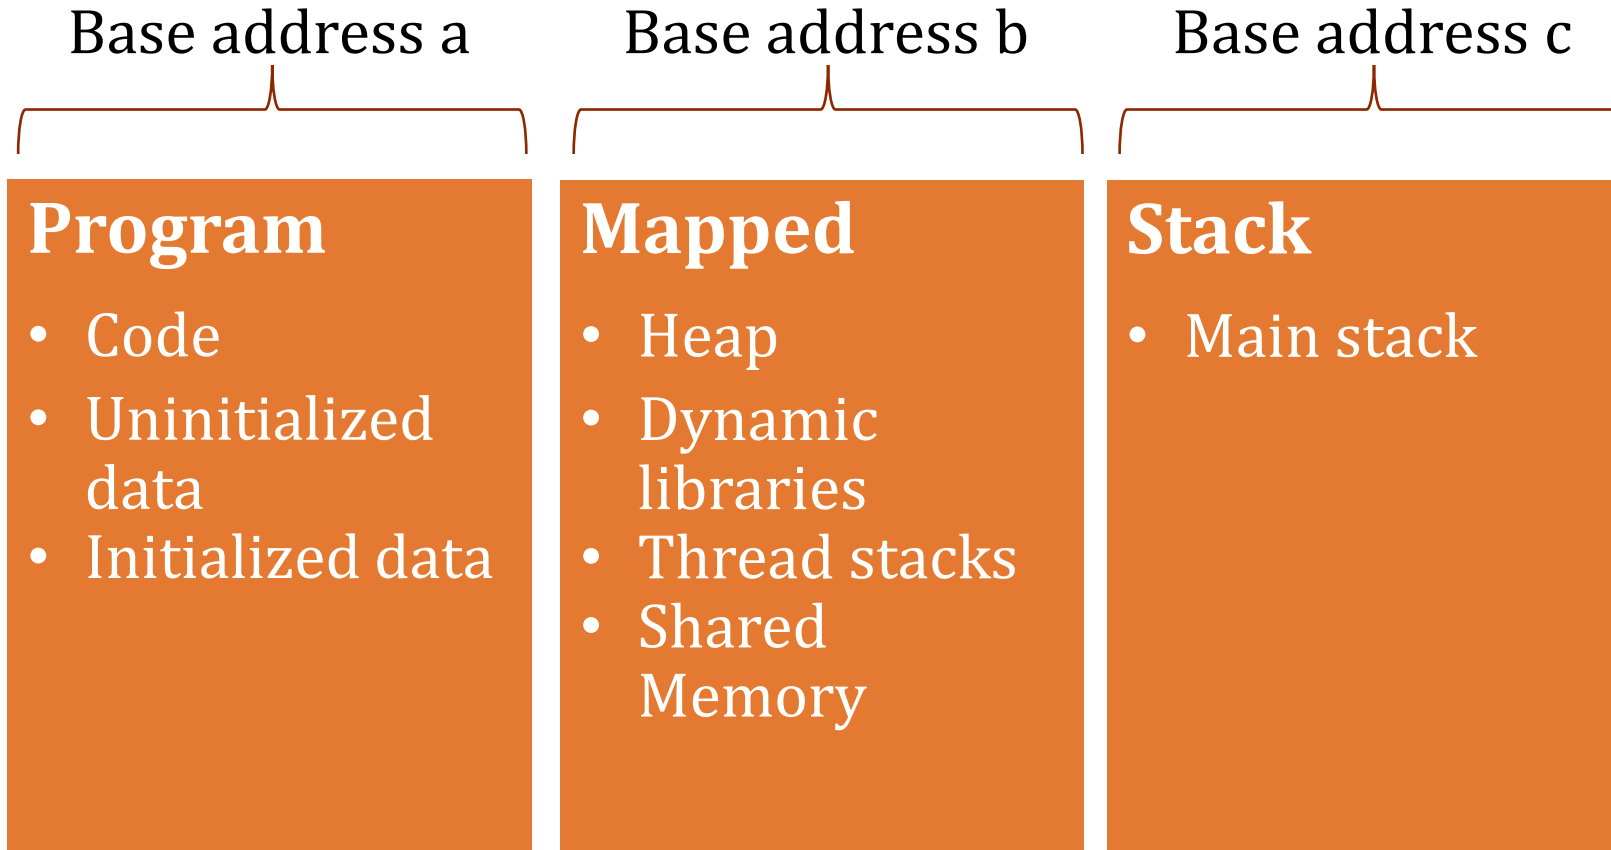
True of both stack-based and return-to-libc:

- Location of shell code
- Library addresses, gadget offsets

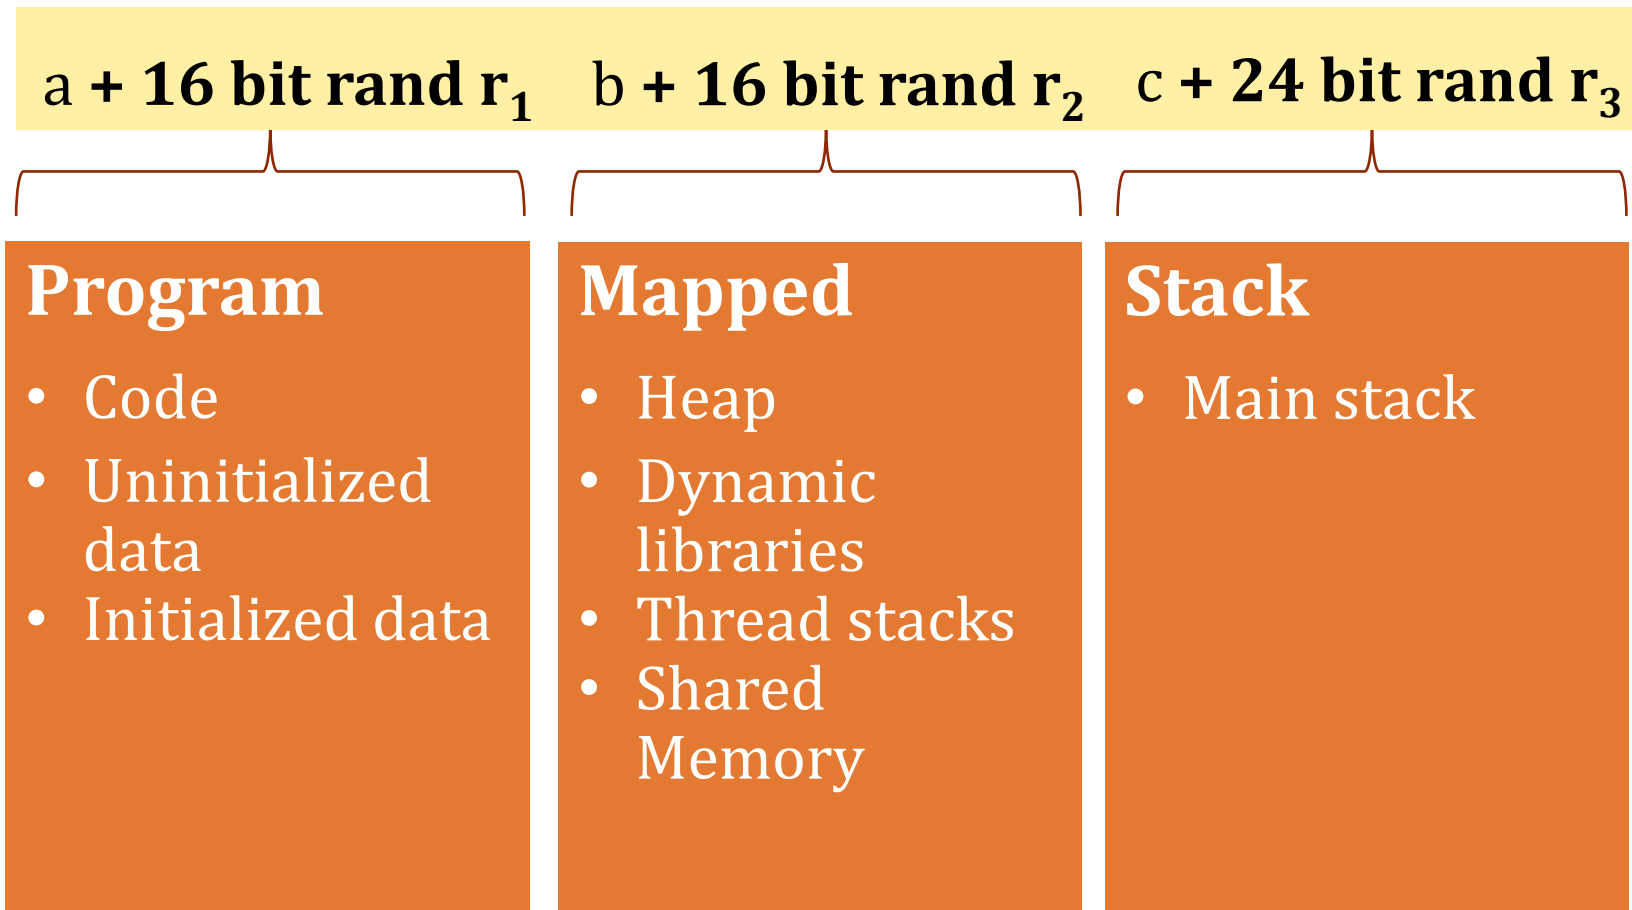
Vulnerability (partially) due to fixed memory layout

**Solution:** Randomize the layout!

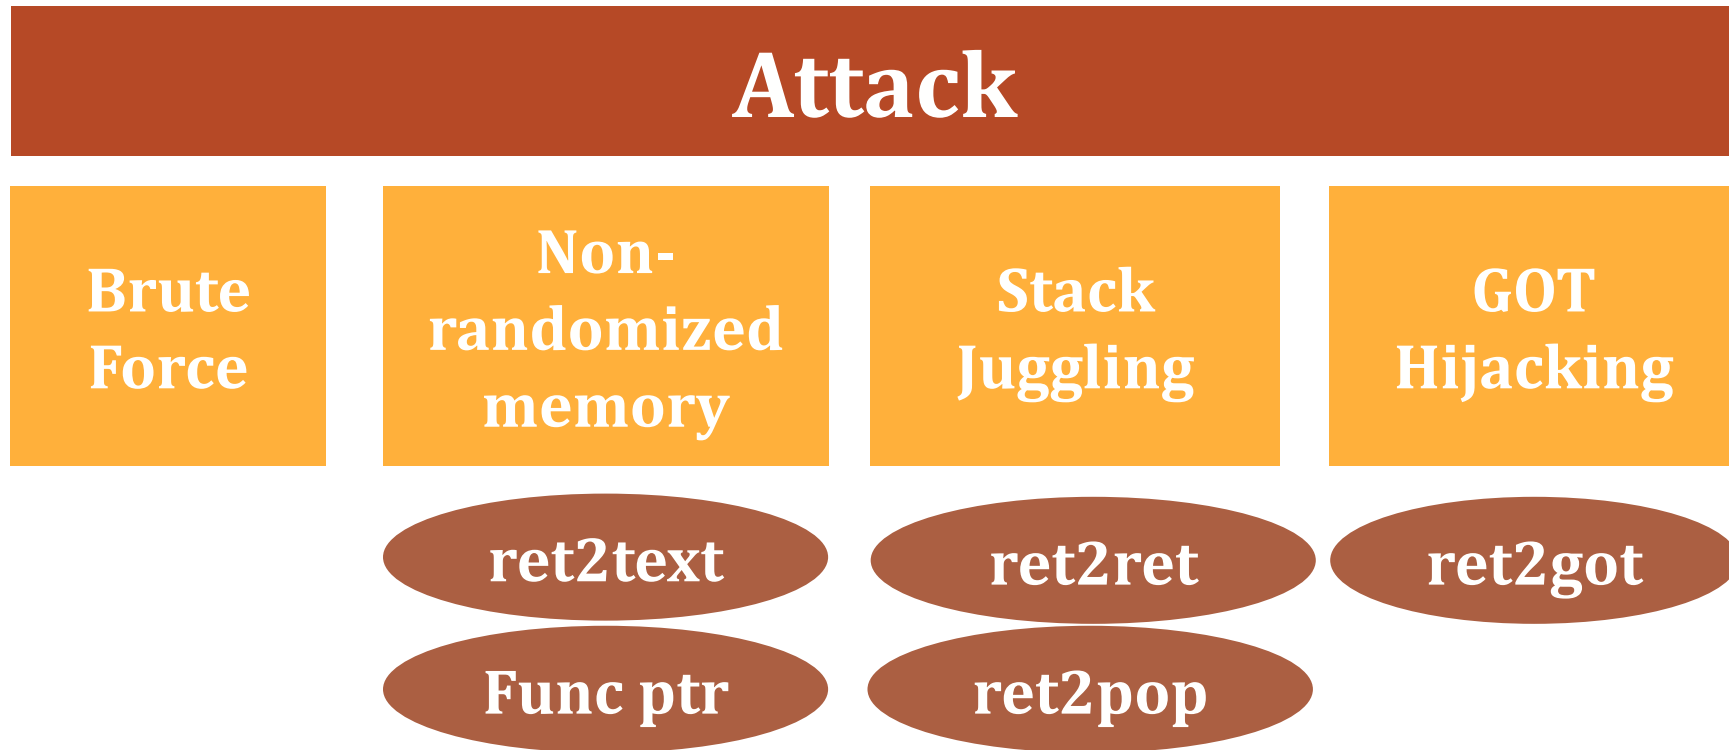
# ASLR



# ASLR



# How to attack ASLR



# Recap

Memory safety (esp. buffer overflow) issues cause vulnerability

There are a wide range of *heuristic* defenses

- Canaries → pointer subterfuge
- Memory protection → return-to-libc
- ASLR → pick your favorite attack

Why not just enforce memory safety?

# Basic Example, revisited

```
int main(int argc, char **argv) {  
    char buf[64];  
    fgets(buf, 64, stdin);  
}
```

How do we know that this  
is actually safe?

How do we know that *this*  
is actually safe?

```
char* fgets(char* s, int n, FILE *iop) {  
    register int c;  
    register char* cs;  
    cs = s;  
    while(--n > 0 && (c = getc(iop)) != EOF)  
    {  
        if((*cs++ = c) == '\n')  
            break;  
    }  
  
    *cs = '\0';  
    return (c == EOF && cs == s) ? NULL : s;  
}
```

# Basic Example, revisited

```
char* fgets(char* s, int n, FILE *iop)
//@requires 0 <= n && 0 < s
{
    register int c;
    register char* cs;
    cs = s;
    while(--n > 0 && (c = getc(iop)) != EOF)
    //@loop_invariant 0 < cs && cs - s <= n
    {
        if((*cs++ = c) == '\n')
            break;
    }

    *cs = '\0';
    return (c == EOF && cs == s) ? NULL : s;
}
```