

Assignment 4: Relaxed Secrecy and Privacy
15-316 Software Foundations of Security and Privacy

Due: **11:59pm**, Thursday 4/12/18

Total Points: 50

1. Safe or unsafe relaxation (15 points).

For the declassification operators defined below, determine whether an attacker can always use it to figure out the value of an n -bit string `pin` in **poly**(n) time. If so, describe how. If not, prove why doing so is impossible using a similar argument to the one for `match` covered in lecture. All declassification operators below have the same typing rule as `match`.

- (5 points) Greater than: `gt(guess, pin)` evaluates `true` if and only if `guess`, interpreted as an integer, is greater than the n -bit string `pin`, interpreted as an integer.
- (5 points) Error-correcting match: `ecm(guess, pin)` evaluates to `true` if and only if `guess` is an n -bit string that differs from the n -bit string `pin` by at most 1 bit.
- (5 points) Prefix: `pref(guess, pin)` evaluates to `true` if and only if `guess` is a prefix of the n -bit string `pin`.

2. Primitive badness (15 points).

RSA is a public key cryptosystem that performs encryption by taking powers modulo N of an exponent e , and decryption by taking powers modulo N of an exponent d . The details of how N , e and d are chosen are not important for this problem, but the pair (e, N) is the *public key* and d is the secret *private key*. To encrypt a plaintext message M , one computes the ciphertext $C = \text{mod}(M^e, N)$. Likewise to perform decryption given C to recover M , one computes $M = \text{mod}(C^d, N)$. Thus modular exponentiation lies at the core of the algorithm, so is the essential primitive needed to implement RSA.

The program below implements modular exponentiation using the square-and-multiply method¹. Given ciphertext C , the approach iterates over each bit j of the L -bit private decryption key d , squaring (mod N) the ciphertext at each step. If the current bit $d[j]$ is 1, then the current result is multiplied by the original ciphertext (again mod N). The modulo operation here is implemented in a very simple manner by repeated subtraction.

```
x := C;
for(j in 0 to L-1) {
  x := x * x;
  while (N <= x) { x := x - N; }
  if (d[j] = 1) {
    x := x * C;
    while (N <= x) { x := x - N; }
  }
}
```

Assuming that all variables except d are public and j is initialized to 0, this modular exponentiation program contains a timing side channel. Explain what it is. Then, given the following timings for each initial value of C below where $L = 4$ and $N = 16$, recover the value of d that led to these observations. You should assume that each arithmetic operation, comparison, and assignment takes one unit of time, and that the `for` loop does not take a unit of time to increment j .

¹You may notice that this code only works when N is relatively prime to C . This is a reasonable assumption for reasons beyond the scope of the assignment, but if you are interested in learning more then we recommend reading *Introduction to Modern Cryptography, Second Edition*, Chapter 3, by Katz and Lindell

| C | $runtime$ | C | $runtime$ | C | $runtime$ | C | $runtime$ |
|-----|-----------|-----|-----------|-----|-----------|-----|-----------|
| 0 | 22 | 4 | 25 | 8 | 30 | 12 | 40 |
| 1 | 22 | 5 | 35 | 9 | 32 | 13 | 52 |
| 2 | 24 | 6 | 29 | 10 | 36 | 14 | 48 |
| 3 | 32 | 7 | 29 | 11 | 46 | 15 | 50 |

3. **Constant-time fix (10 points).** Fix the timing channel in the program from Part 2 so that the runtime no longer depends on the value of d . If it helps make your answer more clear, you can assume that the language contains a $\text{mod}(x, N)$ primitive, but you must also assume that it runs in $\lfloor \frac{x}{N} \rfloor$ units of time. What is the runtime of your fixed implementation?
4. **Randomized enough? (10 points).** Recall the randomized response mechanism discussed in Lecture 14. It flips a fair coin (i.e., one with equal probability 1/2 or returning 0 or 1). If the coin comes heads, then it returned the contents of $\text{Mem}(0)$ (which we assumed to be either 0 or 1). If the coin comes up tails, then it flips another fair coin and returns the value. We saw that this satisfies $\ln(3)$ -differential privacy.

Consider the following variant, which computes a function of both $\text{Mem}(0)$ and $\text{Mem}(1)$.

```

b := flip( $p$ )
if  $b = 1$  then
   $o$  :=  $\text{Mem}(0)$ 
else
   $o$  := flip( $p$ ) +  $\text{Mem}(1)$ 

```

(1)

Use Definition 2 from Lecture 14 to answer this question. Does this program satisfy differential privacy for any value of $\epsilon > 0$? If so, explain why. If not, give a counterexample pair of neighboring databases for which the bound in Equation 8 (Lecture 14) cannot hold for any $\epsilon > 0$, and explain how to modify the program to make it satisfy differential privacy.